


# MatFormer: Nested Transformer for Elastic Inference

Prateek Jain


Google Research India

Devvrit, Sneha Kudugunta, Aditya Kusupati,  
Tim Dettmers, Hannaneh Hajishirzi, Yulia Tsvetkov, Kaifeng Chen, Inderjit Dhillon,  
Sham Kakade, Ali Farhadi

# Large Models: Deployment Story



Deployment Constraints  
(RAM, Latency, QPS...)




PALM 2

- Typically only a few models to choose from
  - Might have to select say Llama 13B even if capacity for Llama 40B
- Distillation/pruning requires additional training


**Goal:** design a “universal” model from which hundreds of accurate models can be extracted

# Existing Solutions towards MatFormer



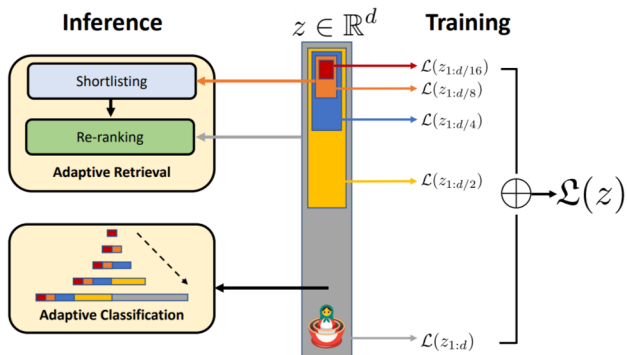
(Universal) Slimmable Networks (Yu & Huang ICLR 2019; ICCV 2019)

- Primarily focused on CNNs
- Training routines w/
  - Modifications to batchnorm
  - Sampled submodels
  - Distillation from largest model
- Longer/costlier training routines

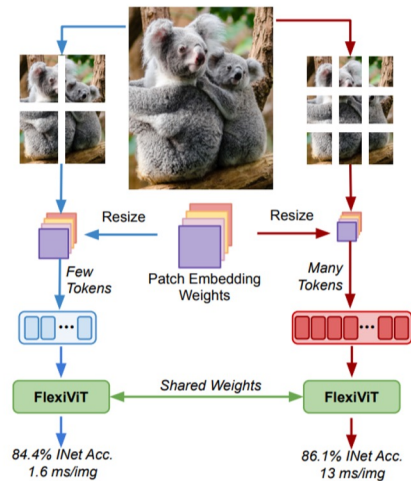


Once-for-All (Cai et al., ICLR 2020)

# Existing Solutions towards MatFormer



Matryoshka Representation Learning (Kusupati et al., NeurIPS 2022)



FlexiViT (Beyer et al., CVPR 2023)


Flexibility in output and input space respectively

# MatFormer: Nested Substructure

- Works for Transformers – without modifications to fundamental blocks
- Joint training: No subsampling or distillation
- Significantly cheaper training cost than equivalent methods while being more accurate
- 4 granularities sufficient to span a wide range of constraints.


# Transformer

Transformer Encoder



**MLP: ~80% of the cost**

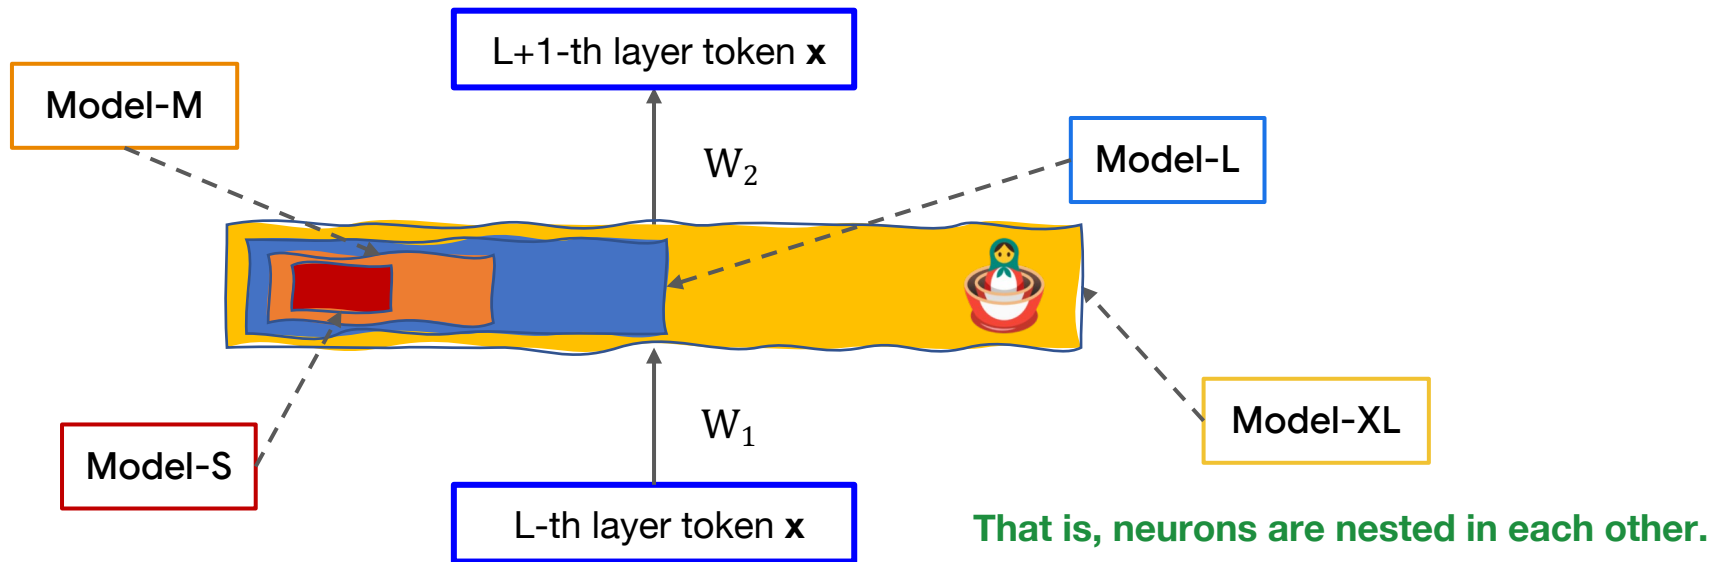
**Attention: ~20% of the cost**



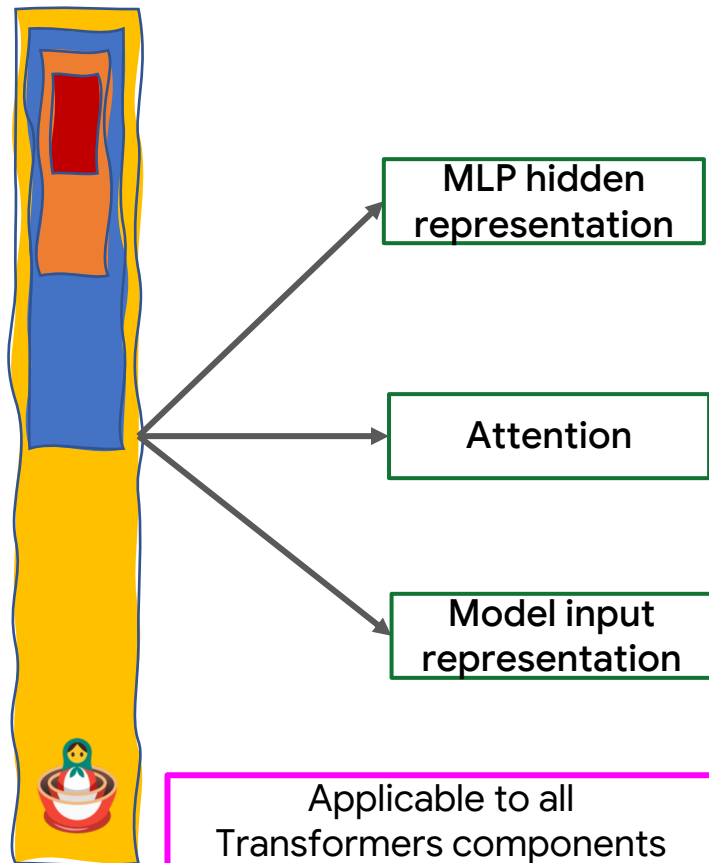
**$h = 4*d$  to  $12*d$  typically**

# MatFormer: Matryoshka Transformer

- **MatFormer** builds upon [MRL](#)
- Apply MRL to MLP layer in each transformer block



# MatFormer: Generality & Training





## Recipe:

- Pick XL model architecture
- Pick  $\mathbf{G}$  granularities for nesting eg.,  $\mathbf{G} = 4$
- Jointly optimize  $\mathbf{G}$  shared models akin to MRL
- Matformer train cost < **total cost of training each granularity from scratch**
- MatFormer can also be induced w/ Fine-tuning



Can generate **1000s** of models not just 4



# Mix'n'Match & Routing on MatFormer



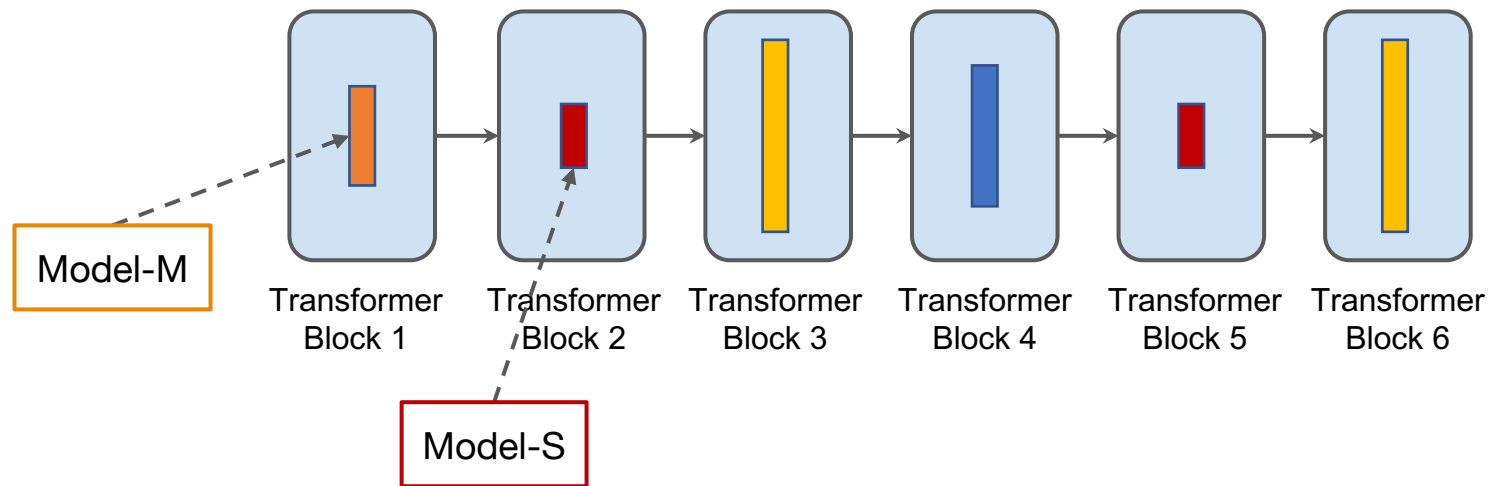
# Mix'n'Match & Routing on MatFormer



This gives only say 4 models.

So where do we get 1000s of models???

# Mix'n'Match & Routing on MatFormer



- **Mix'n'Match:** 100s (combinatorial) of *static (on-demand)* models for all accuracy-compute
- **Routing:** Token based routing akin to MoE to realize *dynamic* computation

# MatLM: MatFormers for Language Modeling

- Standard setting from Lamda (Thoppilan et al.)
- **G = 4** granularities – change the MLP hidden dims!
  - **XL** – hidden\_dim (hd), **L** – hd/2, **M** – hd/4, **S** – hd/8.
- Nomenclature: MatFormer-XL, MatFormer-L, MatFormer-M, MatFormer-S
  - Independently Trained Models: Baseline-XL, Baseline-L, Baseline-M, Baseline-S
- *7 different “XL” model scales: from 78M up to 2.6B parameters.*
  - 78M, 180M, 310M, 463M, 850M, 1.3B, 2.6B

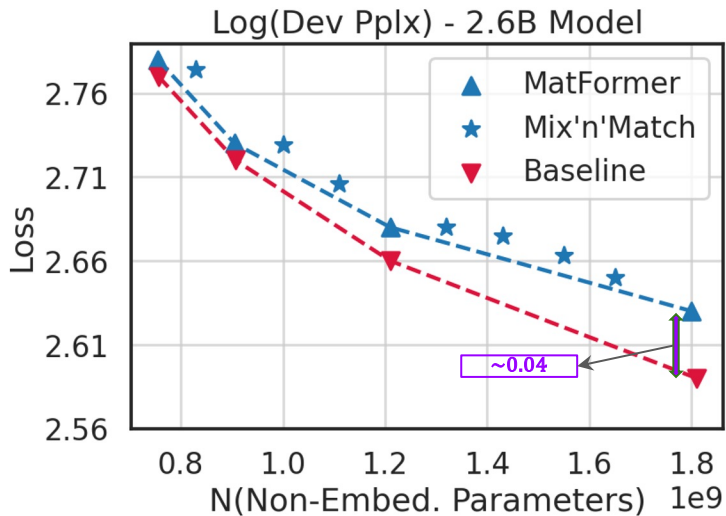
# MatLM: Key Findings

- Little to **no loss in test pplx** and **GPT3 1-shot downstream evals**.
  - For each granularity, i.e.,  $\text{accuracy}(\text{MatFormer-Z}) \sim \text{accuracy}(\text{Baseline-Z})$
  - $Z \in [\text{XL}, \text{L}, \text{M}, \text{S}]$
- Able to read models for *free* using Mix'n'Match
  - Mix'n'Match interpolates well between the 4 granularities
- Side effects: consistency with large model, gains over Speculative Decoding

# Language Modeling with 2.6B model: Mix'n'Match

Trained for: 160B tokens

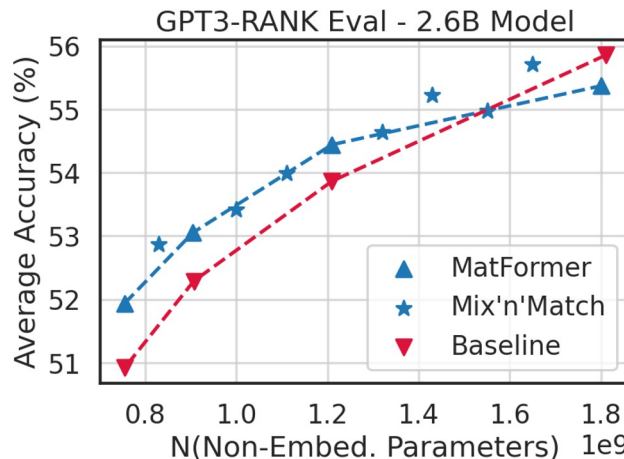

Decoder-only model w/ same hparameters for baseline & MatFormer



- XL model size: 2.6B
  - We read-off L, M, S models for matformer
  - Baselines trained from scratch for all granularities
- Log pplx within 0.04 of the baseline for 2.6B model!
  - Downstream evals almost match (see next slide)
- All the ★ are for “free” – *they were never trained for*
  - We just read them from Matformer-XL

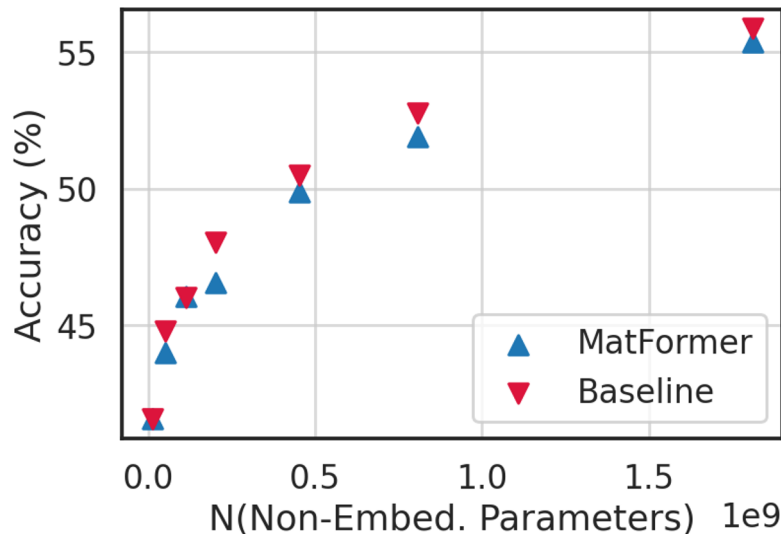

# Language Modeling with 2.6B model: Mix'n'Match

1-shot GPT-evals



- Almost matching accuracy for MatFormer-[XL, L, M, S] models against Baselines
- We get all the intermediate models denoted by ★ for “free”
  - No extra training!
  - For GPT-3 Rank: ★ models almost lie on a line interpolating trained MatFormer models (XL, L, M, S)

# Language Modeling: Scaling Plots for Universal Model (XL)



**MatFormer** :  $23.0528 * N^{-0.1407} + 1 / D + 1.4352$

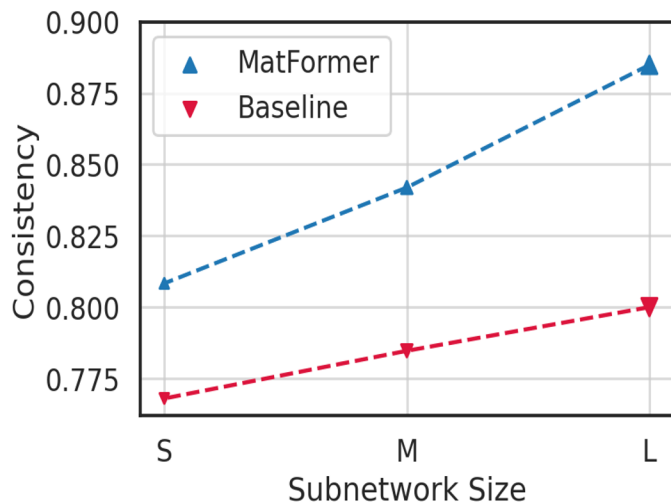
**Baseline** :  $22.7879 * N^{-0.1414} + 1 / D + 1.4973$

Where N = N(Non-Embedding Parameters) and D = N(Training Tokens)

Above laws hold for all trained granularities: XL, L, M, S!



# Language Modeling: Consistency for 2.6B XL model



**Consistency:** accuracy of smaller models (S, M, L) when output of XL model is the ground truth

## ***Why care about consistency?***

*Techniques like Speculative Decoding becomes more efficient with more consistent models*

Speculative Decoding	LAMBADA	TriviaQA
Baseline	1.10×	1.08×
MatLM	1.14×	1.11×
+ shared attention cache	1.16×	1.14×


*MatFormer subnetworks are significantly more consistent with the full model compared to vanilla baselines.*

# MatViT: MatFormer + ViT


- Generalized formulation translating to ViT
- Works for across model sizes for both pre-training and fine-tuning
- Enables accurate adaptive encoders for classification
  - Spans all of the space with Mix'n'Match (and potentially routing)
- Enables accurate adaptive query encoders for retrieval
  - Use the largest model for Index building
  - Leverage smaller query encoders during inference based on the constraints
  - This requires aligned training/distillation for baseline models to work

# MatViT: Classification

ViT-B/16 -- ImageNet-1K



ViT-L/16 -- ImageNet-1K (PT IN-21K)

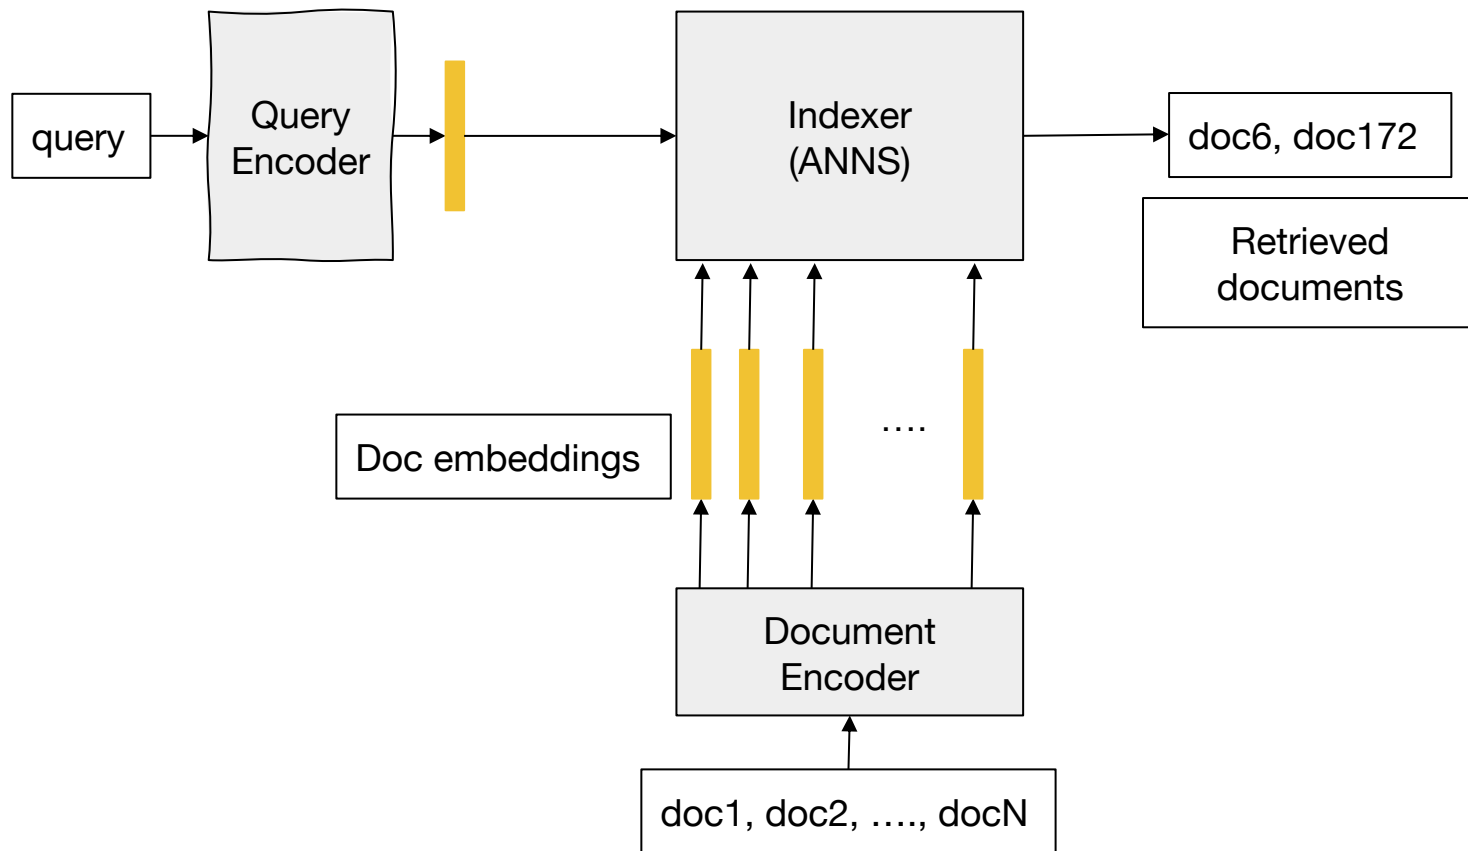


All the ★ are for “free” during inference – *they were never optimized for.*

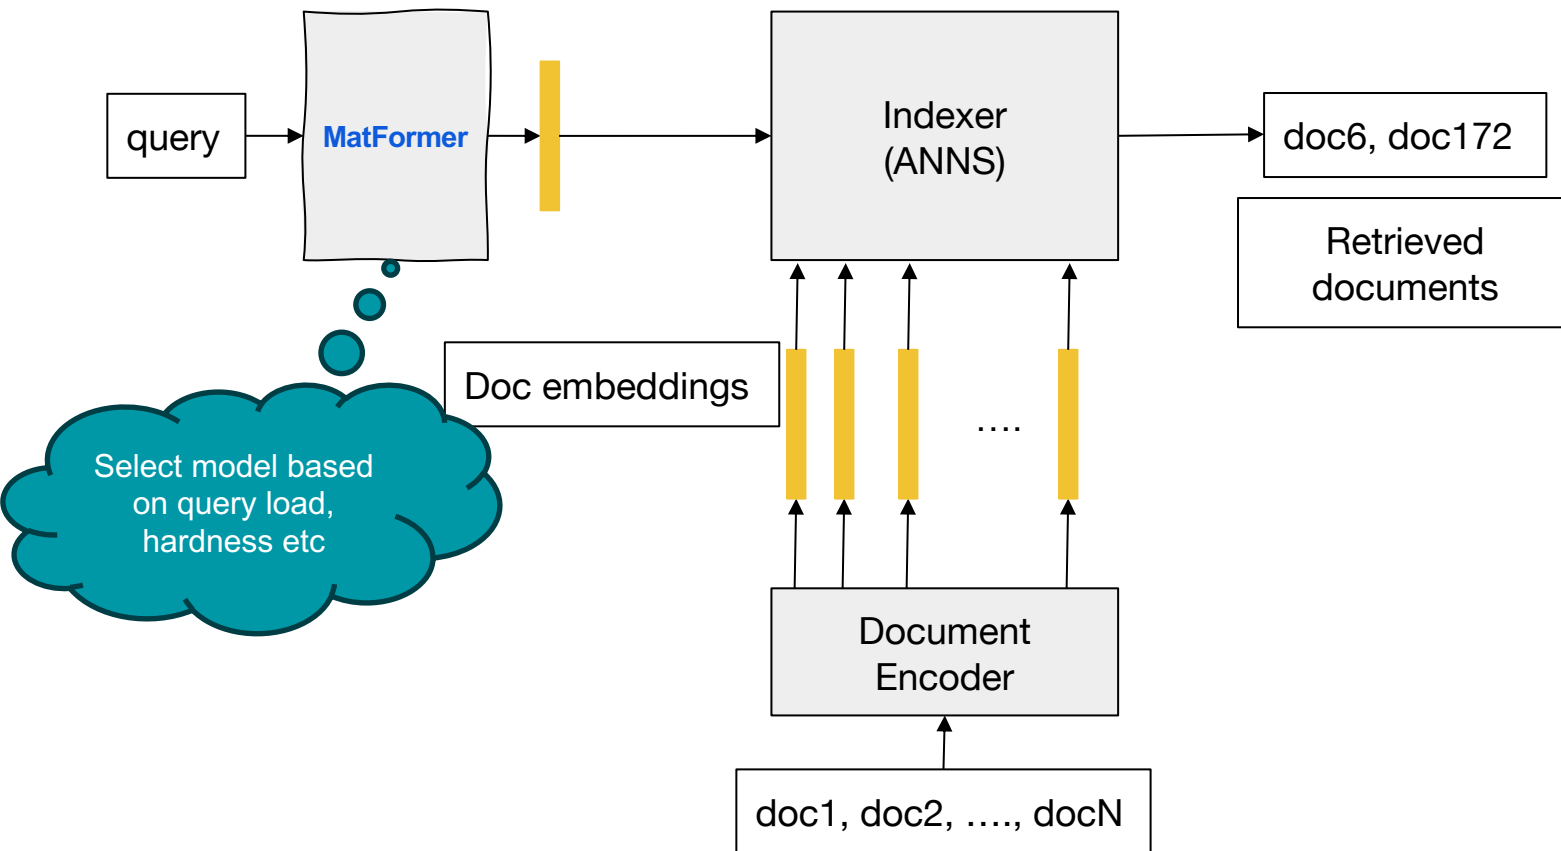
# MatViT: MatFormers + ViT

- Generalized formulation translating to ViT
- Works for across model sizes for both pre-training and fine-tuning
- Enables accurate adaptive encoders for classification
  - Spans all of the space with Mix'n'Match (and potentially routing)
- Enables accurate adaptive query encoders for retrieval
  - Use the largest model for Index building
  - Leverage smaller query encoders during inference based on the constraints
  - This requires aligned training/distillation for baseline models to work



# Semantic Search: Dual Encoder Models



# Semantic Search: Flexible Dual Encoder Model



# MatViT: Adaptive Retrieval (Index built w/ largest model)



All the ★ are for “free” during inference & preserve metric space.

# MatFormer + ViT-B/16: Cross-consistent Retrieval

1-NN accuracy (%) with varying index and query encoder sizes from MatViT-B/16  
(Baseline numbers): Rest are near Random

Index↓/Query→	36M	43M	57M	85M
36M	<b>72.42%</b> (71.44%)	74.31%	75.33%	76.26%
43M	72.30%	<b>74.71%</b> (74.90%)	75.93%	76.69%
57M	72.12%	74.71%	<b>76.44%</b> (76.58%)	77.19%
85M	71.71%	74.48%	76.40%	<b>77.40%</b> (77.46%)

Adaptive Query Encoders for Retrieval

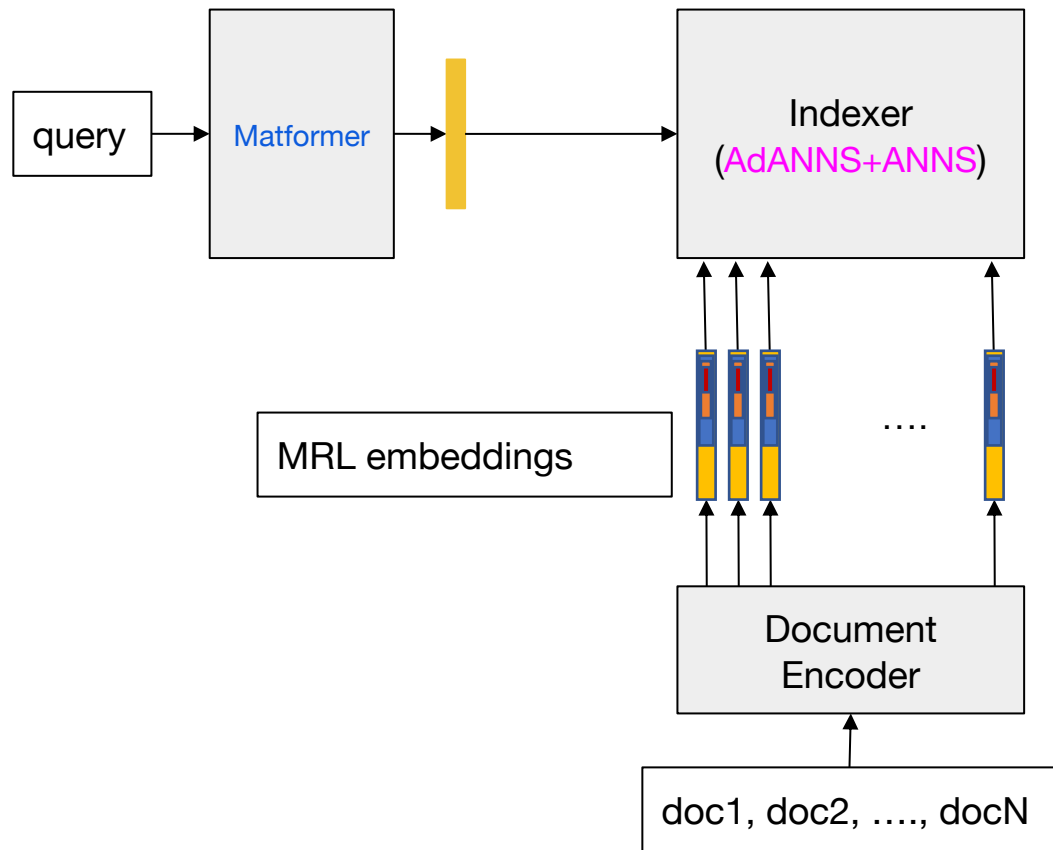


# Summary and Future Work


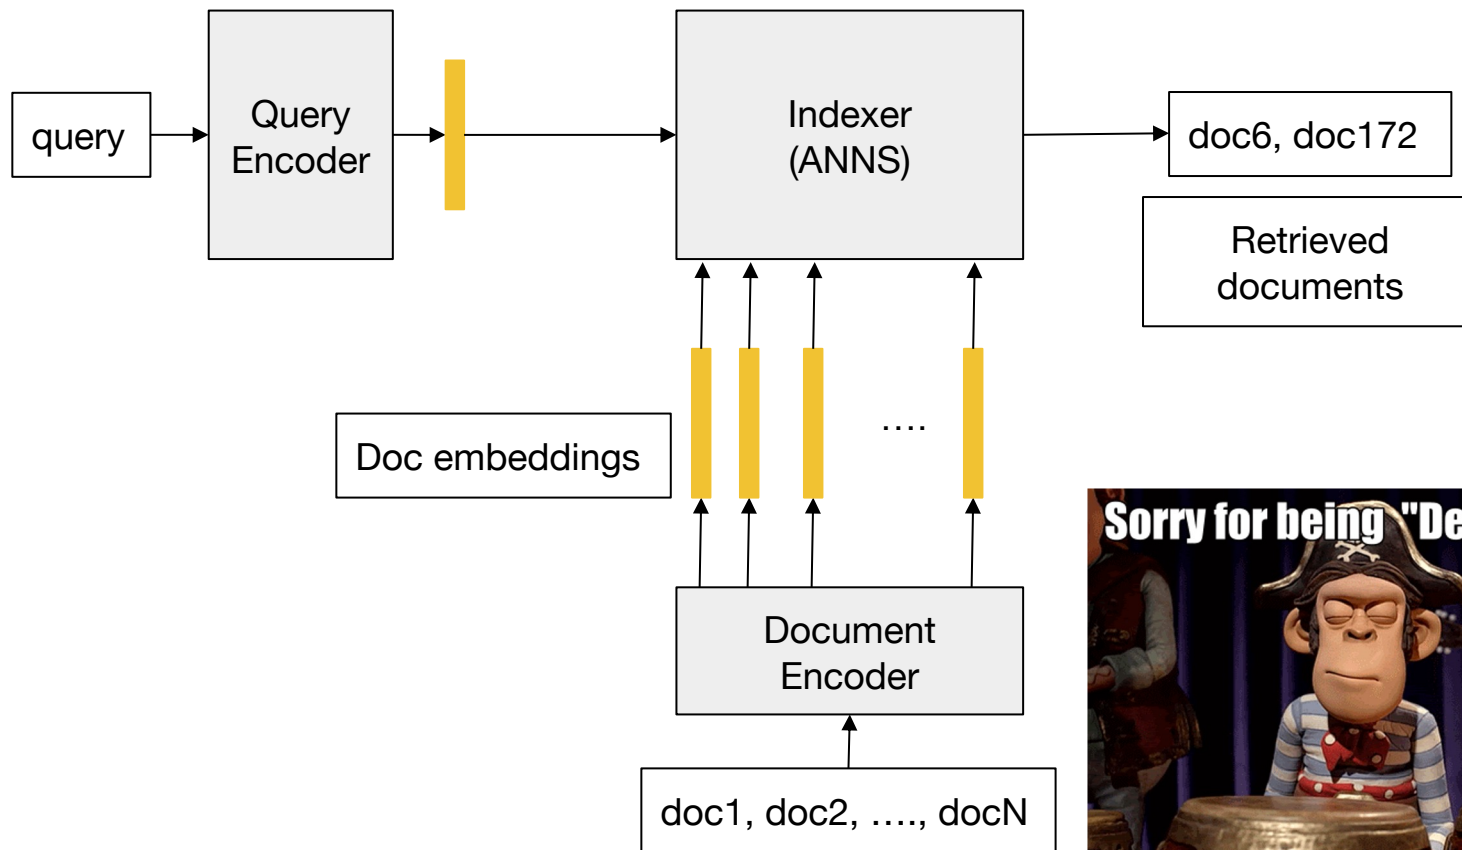
- Matformer: nested substructure for elastic inference
- Training: joint optimization of a few granularities
- MatLM: Matformer Language Model
  - 2.6B scale models with same pplx and evals as independently trained baselines
  - Consistency: gains over speculative decoding
- MatViT: Matformer Vision Transformer
  - ViT-L/16 scale models, similar performance as independently trained baselines
  - Adaptive retrieval

## Future Work:

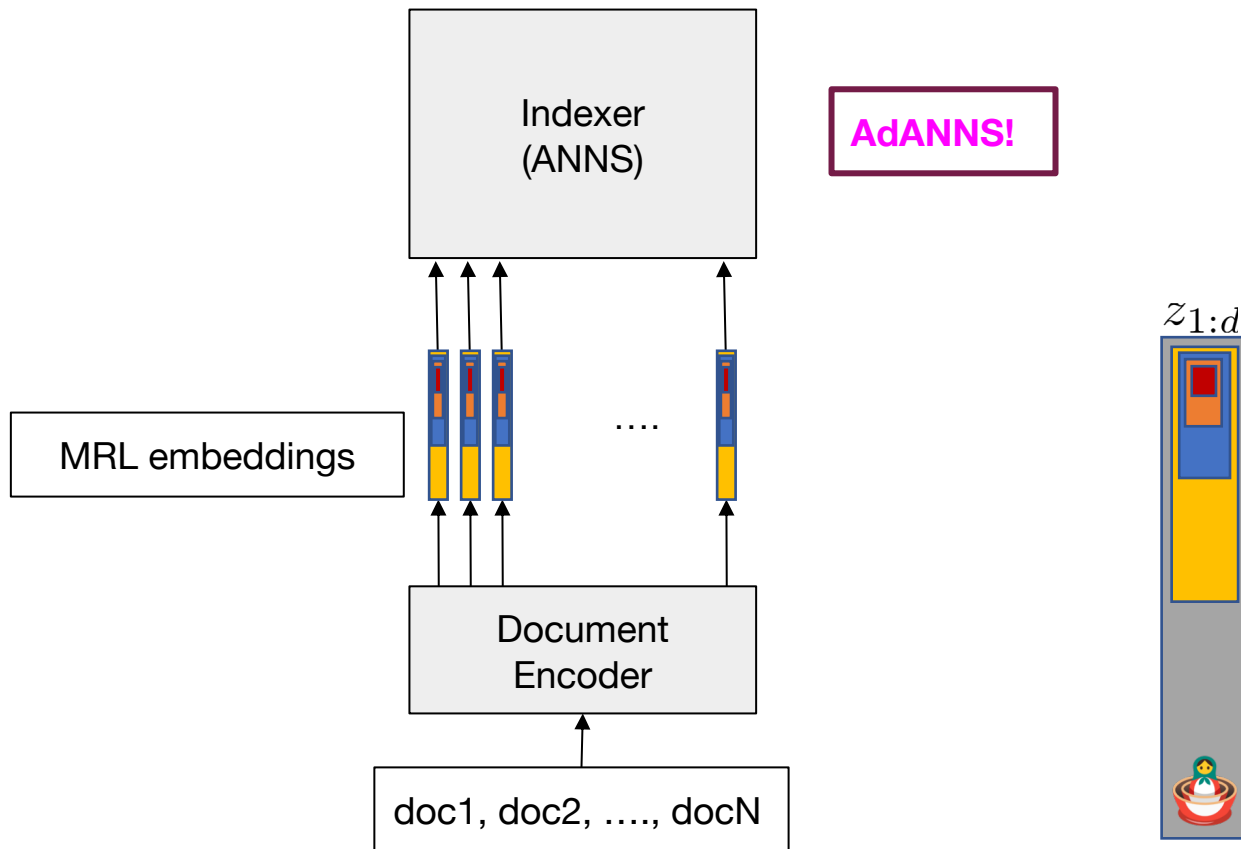
- Further investigation of scaling laws
- Better training algorithms
- Practical deployment of a truly elastic system



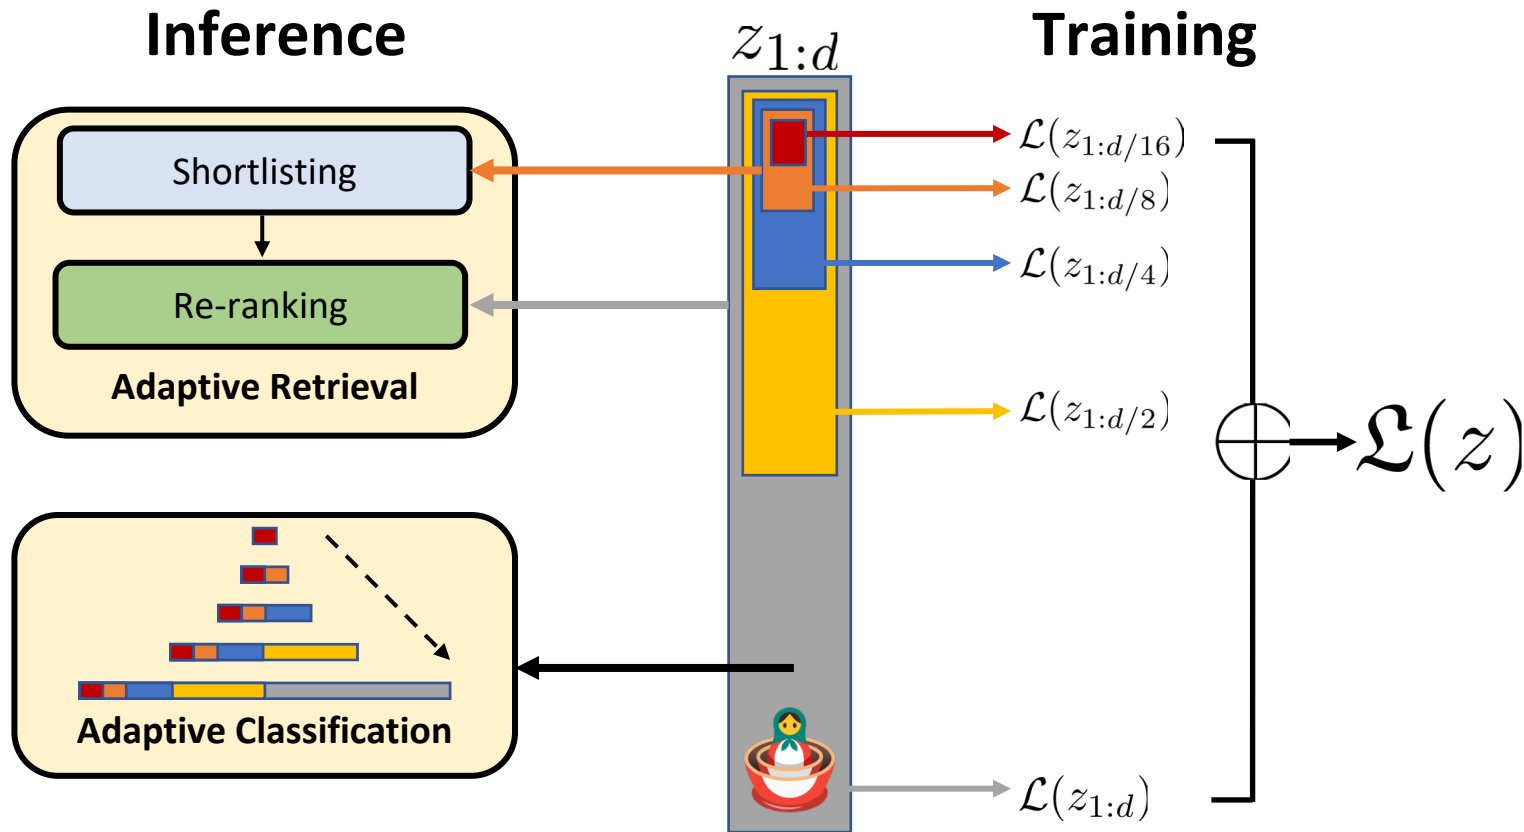
# Semantic Search via Dense Retrieval




# Mission: Make entire dense retrieval system “adaptive”



# Matryoshka Representation Learning - MRL




# ANNS-IVF



- Clustering through **k-means**
- Can vary number of clusters and leaves chosen
- Within each leaf node – Linear Scan

**Regular IVF**

# Adaptive Retrieval with MRL + ANNS (AdANNS)




- **More accurate** at same cost
- **~1.25x cheaper** for same accuracy
- More flexibility in design


<https://arxiv.org/abs/2305.19435>

**AdANNS-IVF**

# AdANNS vs Status Quo (Clustering and Quantization)



(a) Image retrieval on ImageNet-1K.



(b) Passage retrieval on Natural Questions.



